



Qwen-2.5-7B shapes · A10G · BF16 · cache budget 1,024


One layer · 131,072 tokens




One layer · additional allocation



All 28 layers · 32,768 tokens



All 28 layers · absolute allocation



Whole model: common MLP chunk size 1,024; default allocator; excludes LM head and decode.
All tested whole-model configurations OOM at 131,072 tokens. Quality is evaluated separately.